Analyzing the Workout Dataset Using Data Science Techniques

(3130 words, ~23 mins reading time)

Author: Upama Thapa Magar

Introduction: With increasing fitness programs, health apps, and data insights becoming more popular, people are being more interested in understanding how calories are burned and the science behind it. Tracking calories burned is more than just a metric—it's an essential tool for individuals aiming to reach their fitness goals, maintain a balanced diet, or optimize their workout routines. However, the factors influencing caloric expenditure go beyond simple metrics like distance walked or steps taken. Elements such as workout type, session duration, age, weight, and even exercise experience can significantly impact how many calories are burned during a workout. In this blog, I will explore a dataset of gym members that contains details about workout habits and individual characteristics. By using machine learning, I aim to identify how demographic factors, exercise frequency, and other variables affect calorie burn. Through data analysis and model building, I'll uncover trends and develop a predictive model that estimates caloric burn based on a combination of factors, including workout type, experience level, and more. This project not only provides insights into fitness behavior but also highlights the power of data science in creating personalized health recommendations. With these tools, we can better understand what drives effective workouts and, ultimately, improve the way we approach fitness.

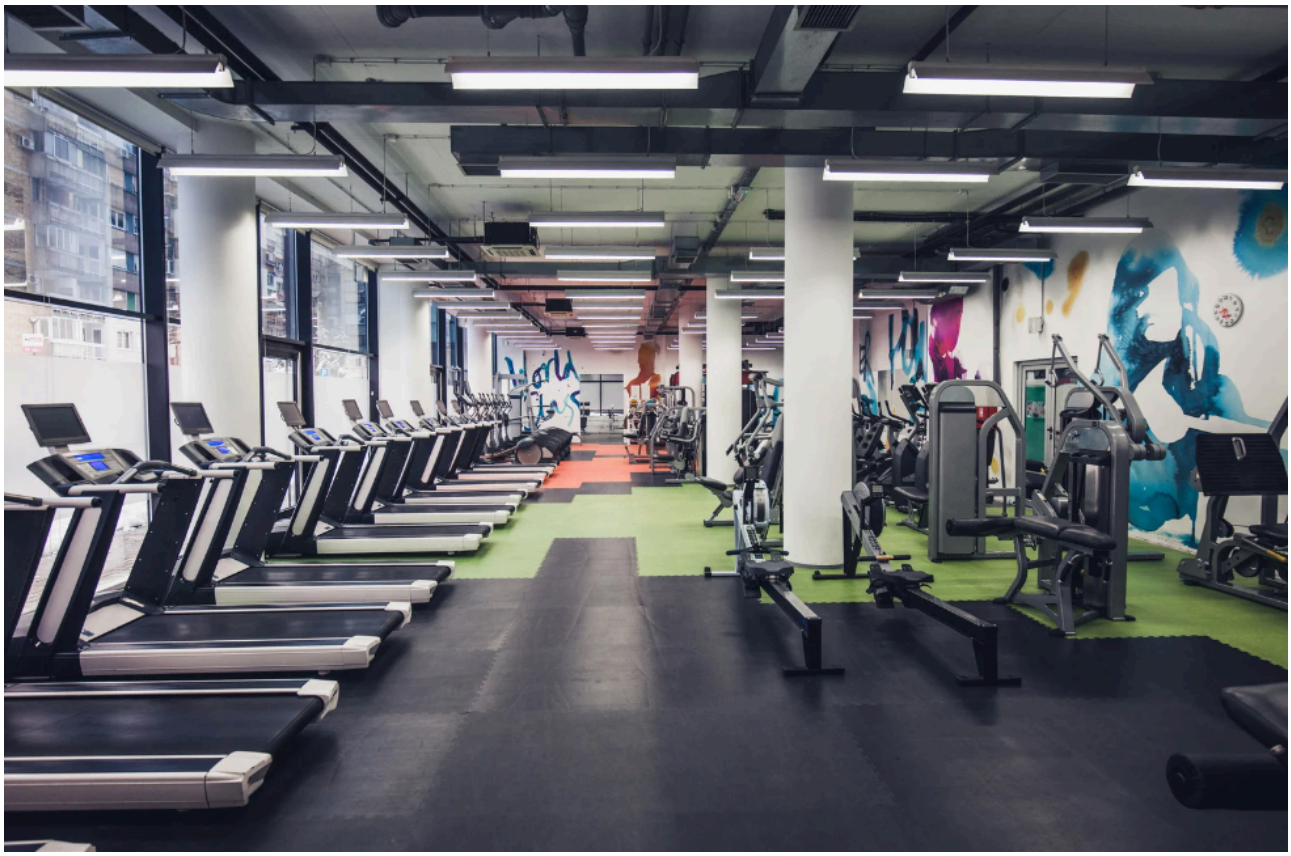

How can we uncover insights and make predictions from gym workout data?

This project explores the vast potential of gym workout data to provide insights into fitness habits and predict outcomes like calories burned, workout effectiveness, and user engagement. By analyzing various factors, such as workout frequency, experience level, average heart rate, and fat percentage, I aim to understand how these variables contribute to energy expenditure and fitness outcomes. Through exploratory data analysis (EDA), I identified key patterns and relationships, then applied machine learning models—such as Linear Regression, Decision Tree, and Random Forest—to make predictions based on these features. Evaluating each model's performance with metrics like Mean Squared Error (MSE) and R2 score, and assessing feature importance, provided a clear view of which factors impact outcomes most significantly. This comprehensive approach not only helps in making accurate predictions but also offers actionable insights for personalized fitness recommendations, with detailed explanations in the notebook for a deeper understanding of each step and result.

Data Overview: For this analysis, I used a detailed dataset from Kaggle's Gym Members Exercise Dataset, which provides insights into various characteristics and workout behaviors of gym members. This dataset is rich with both demographic and workout-related information, offering a comprehensive look at how different factors might influence caloric burn during exercise sessions. Here is a breakdown of the key features in the dataset:

- **Age**: The age of each gym member. This demographic variable may influence workout habits and caloric burn, as age can affect metabolism, endurance, and workout preferences.
- **Gender**: Gender of each member, recorded as "Male" or "Female." Gender can be an important factor in fitness analytics since physiological differences may lead to variations in workout routines, intensity, and caloric expenditure.
- **Weight (kg)**: Each member's weight in kilograms. Weight is a crucial variable for caloric burn as individuals with higher body weights generally burn more calories during physical activities due to increased energy requirements.
- **Height (m)**: The height of each member, recorded in meters. Along with weight, height can be used to calculate Body Mass Index (BMI), which is included as a separate feature in this dataset.
- **Max_BPM**: The maximum beats per minute (BPM) recorded during a workout session. This feature could provide insights into the intensity of workouts, as higher BPM often correlates with more strenuous activities.
- **Avg_BPM**: The average beats per minute throughout a session. Like Max_BPM, Avg_BPM reflects workout intensity, and tracking both can reveal variations in cardiovascular strain between different types of exercise.
- **Resting_BPM**: The resting heart rate of each member before starting their workout. Resting BPM serves as a baseline to compare against Max_BPM and Avg_BPM, offering a clearer picture of workout intensity relative to individual fitness levels.
- **Session_Duration (hours)**: The duration of each workout session, measured in hours. Longer sessions might indicate endurance-focused workouts, while shorter sessions could point to high-intensity interval training (HIIT) or strength-based routines.
- **Calories_Burned**: This is the target variable representing the total calories burned during each workout session. Understanding the relationship between this variable and other features is the primary focus of this analysis, as it directly relates to workout effectiveness and fitness goals.

- **Workout_Type**: The type of workout performed, categorized as "Cardio," "Strength," "HIIT," or "Yoga." Each type of exercise has different characteristics and expected energy expenditures, making it a valuable feature for analyzing caloric burn patterns.
- **Fat_Percentage**: The body fat percentage of each gym member. Fat percentage is an indicator of body composition, which can affect metabolism and caloric needs, thus influencing how many calories are burned during different types of workouts.
- **Water_Intake (liters)**: The amount of water consumed by each member in liters. Hydration is essential for optimal physical performance and recovery, and adequate water intake can influence workout duration and intensity.
- **Workout_Frequency (days/week)**: The number of days per week each member exercises. Workout frequency reflects workout habits and dedication, and frequent exercise can lead to improved fitness levels, potentially impacting caloric burn and overall energy expenditure.
- **Experience_Level**: The member's level of experience with workouts, categorized as beginner, intermediate, or expert. This feature may indicate the member's familiarity with exercises and their ability to handle intense or complex workout routines.
- **BMI**: Body Mass Index, calculated based on weight and height. BMI is a widely used indicator of body composition and can provide insights into each individual's health status, which might indirectly affect their workout efficiency and caloric burn rates.

Each of these features plays a role in shaping the workout profiles of gym members and provides a unique perspective on fitness habits and caloric expenditure. By analyzing these features together, I aim to build a deeper understanding of how various factors correlate with the calories burned during exercise sessions, offering insights into optimizing workouts for better results.

Data Preprocessing: In any data science project, preprocessing is a crucial stage that prepares the raw data for analysis and modeling. For this dataset, I employed several preprocessing techniques to ensure the data was clean, consistent, and structured for machine learning.

Note: Detailed explanations of each figure, analysis process, and interpretation of the results are included in the notebook for a comprehensive understanding of the methods and insights gained.

1. **Initial Exploration and Assessment**
   To understand the dataset's structure and quality, I conducted an initial assessment by examining summary statistics and general information about the features. Using functions like .info() and .describe(), I gained insights into the data types, potential missing values, and basic statistical properties (e.g., mean, median, and standard deviation). This initial exploration helped identify areas needing attention, such as missing values or inconsistent data types, guiding the rest of the preprocessing steps.

	Age	Gender	Weight (kg)	Height (m)	Max_BPM	Avg_BPM	Resting_BPM	Session_Duration (hours)	Calories_Burned	Workout_Type	Fat_Percentage	Water_Intake (liters)	Workout_Fre (days
0	56	Male	88.3	1.71	180	157	60	1.69	1313	Yoga	12.6	3.5	
1	46	Female	74.9	1.53	179	151	66	1.30	883	HIIT	33.9	2.1	
2	32	Female	68.1	1.66	167	122	54	1.11	677	Cardio	33.4	2.3	
3	25	Male	53.2	1.70	190	164	56	0.59	532	Strength	28.8	2.1	
4	38	Male	46.1	1.79	188	158	68	0.64	556	Strength	29.2	2.8	
...	...	...	...	...	...	...	...	...	...	...	...	...	
968	24	Male	87.1	1.74	187	158	67	1.57	1364	Strength	10.0	3.5	
969	25	Male	66.6	1.61	184	166	56	1.38	1260	Strength	25.0	3.0	
970	59	Female	60.4	1.76	194	120	53	1.72	929	Cardio	18.8	2.7	
971	32	Male	126.4	1.83	198	146	62	1.10	883	HIIT	28.2	2.1	
972	46	Male	88.7	1.63	166	146	66	0.75	542	Strength	28.8	3.5	

973 rows × 15 columns

2. **Handling Missing Values**
   Missing data can significantly impact model performance and interpretability, so it was essential to address any gaps. I began by calculating the total missing values per feature to understand their distribution across the dataset. Depending on the feature's context, I handled missing values through imputation (e.g., filling with the median or mode) or by removing rows with substantial missing values. This step was critical for maintaining the dataset's integrity and ensuring that all features contribute reliably to model training.
3. **Encoding Categorical Variables**
   Many machine learning algorithms require numerical input, making it necessary to convert categorical variables into a compatible format. I used label encoding to transform categorical features, such as Gender, Workout_Type, and Experience_Level, into numerical representations. Label encoding assigns a unique integer to each category within a feature, which allows the model to interpret these attributes while preserving their categorical nature. This conversion was essential for incorporating categorical insights without sacrificing the model's numerical requirements.

```
Age                            0
Gender                         0
Weight (kg)                    0
Height (m)                     0
Max_BPM                        0
Avg_BPM                        0
Resting_BPM                    0
Session_Duration (hours)       0
Calories_Burned                0
Workout_Type                   0
Fat_Percentage                 0
Water_Intake (liters)          0
Workout_Frequency (days/week)  0
Experience_Level               0
BMI                            0
dtype: int64
```

4. **Scaling Numerical Features**
   In datasets where numerical features exhibit varying scales, larger-value features can

disproportionately influence model training. To address this, I applied standard scaling, normalizing each numerical feature (e.g., Age, Weight, Height, Session_Duration) to have a mean of zero and a standard deviation of one. Standard scaling ensures that all numerical features contribute equally to the model, preventing those with larger ranges from overshadowing smaller-scale variables. This step is particularly beneficial for algorithms that rely on distance measurements, such as K-Nearest Neighbors.

5. **Outlier Detection and Removal**
   Outliers are extreme values that can skew model predictions and obscure general patterns in the data. I employed the Interquartile Range (IQR) method to identify and remove outliers across numerical features. The IQR method calculates the spread between the first (Q1) and third quartiles (Q3), defining outliers as values falling below Q1 - 1.5 * IQR or above Q3 + 1.5 * IQR. Removing these outliers helped refine the dataset, making it more representative of typical gym member behaviors and workout patterns, which in turn enhanced model reliability.
6. **Train-Test Split**
   To evaluate model performance objectively, I divided the data into training and testing sets using an 80-20 split ratio. The training set (80%) allowed the model to learn patterns in the data, while the testing set (20%) provided an independent sample to assess the model's predictive accuracy. This division is crucial for understanding the model's generalization ability, as it reveals how well the model can handle unseen data. By using a separate testing set, I ensured that performance metrics reflect real-world applicability rather than overfitting to the training data.

Exploratory Data Analysis (EDA): During EDA, I visualized the relationships between key variables. For this dataset of gym members, EDA provided insights into how different factors influence caloric expenditure, workout behaviors, and overall fitness patterns. Below are some of the key analyses and visualizations I conducted.

1. **Distribution of Demographic and Workout Features**
   To start, I examined the distribution of demographic variables such as Age, Gender, and Experience Level. Visualizing age distribution, for instance, helped reveal the predominant age groups in the dataset, while analyzing Gender and Experience Level distributions showed the balance across different groups. These distributions are essential for understanding the general characteristics of the gym members and identifying any imbalances that might affect model predictions.

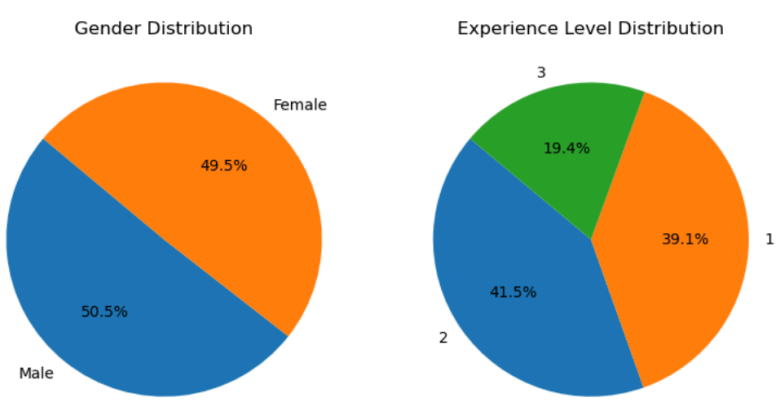

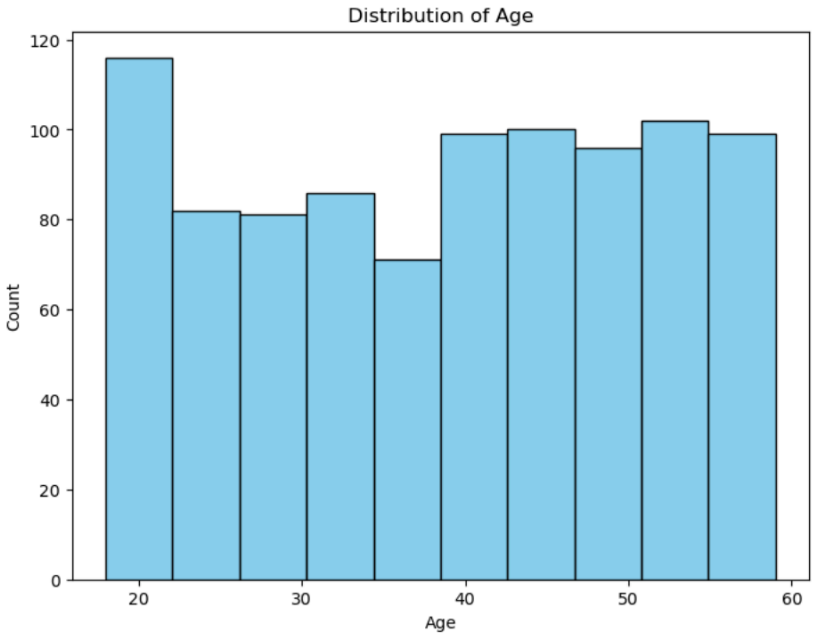

2. **Workout Type Popularity**

Understanding the popularity of different workout types was another area of focus. Using bar charts, I visualized the count of each workout type (e.g., cardio, strength training, HIIT, and yoga) across the dataset. This analysis not only highlighted the most common workouts but also provided a baseline for further comparisons, such as how workout type relates to calories burned or session duration.

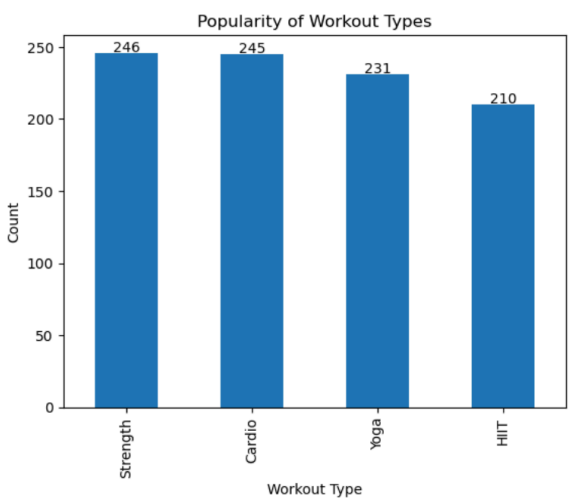

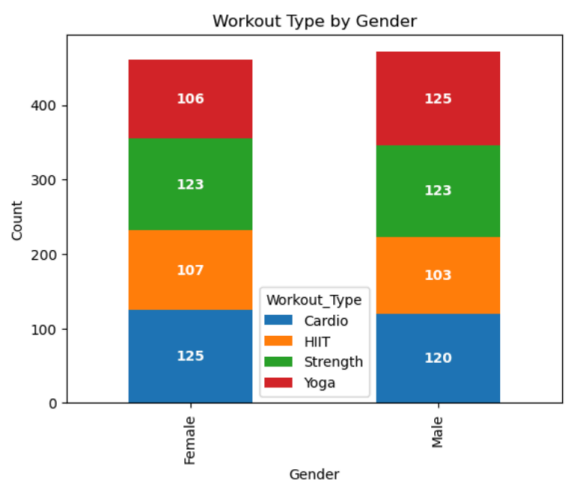

3. **Relationship between Workout Type and Calories Burned**

Calories burned is a core target variable in this dataset, so examining its relationship with other features was crucial. For instance, I used box plots to visualize the caloric expenditure across different workout types. This comparison allowed me to observe the variation in calories burned

based on workout type, showing, for instance, whether HIIT workouts or strength training burned more calories on average. This insight is essential for interpreting how exercise choices impact energy expenditure.

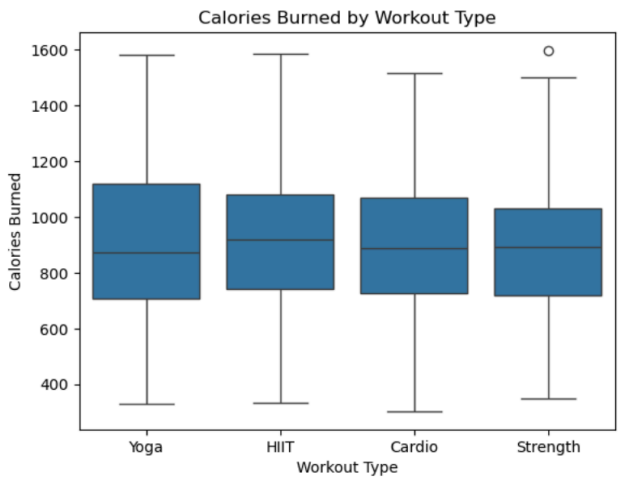

4. **Impact of Demographics on Workout Patterns**
   I also analyzed how demographic factors, such as Gender and Experience Level, influenced workout behaviors. For example, I plotted the average session duration and workout frequency by gender and experience level, using grouped bar charts. These visualizations helped uncover patterns, such as whether more experienced gym members worked out more frequently or whether there were notable differences in session durations across genders. Such insights help contextualize workout preferences and behaviors.

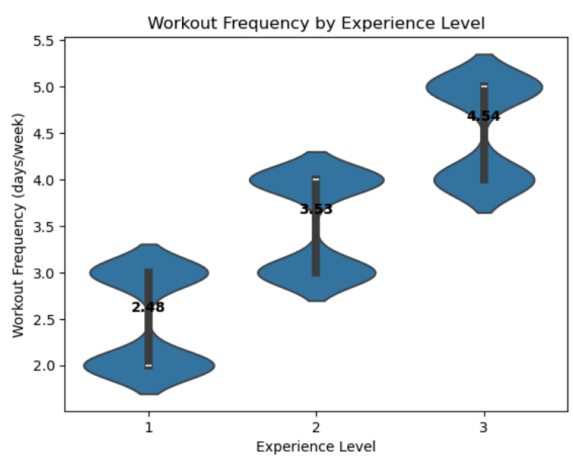

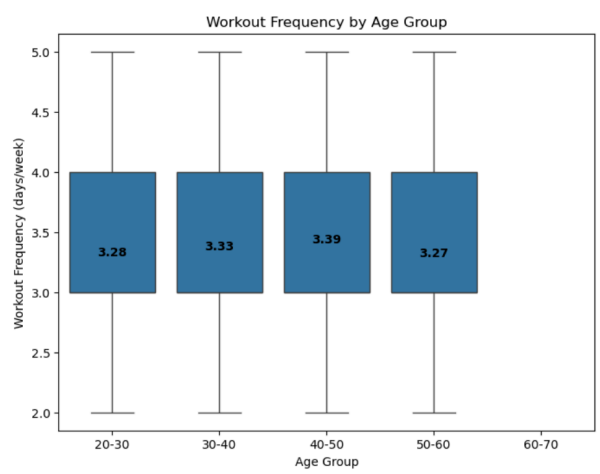

5. **Correlation Analysis of Continuous Variables**
   To identify relationships between continuous variables, I conducted a correlation analysis and visualized the results with a heatmap. Features such as Calories Burned, BMI, Session Duration, Water Intake, and heart rate metrics (e.g., Max_BPM, Avg_BPM) were analyzed for correlations.

The heatmap provided an at-a-glance understanding of how these features are interrelated. For instance, a strong positive correlation between Session Duration and Calories Burned highlights the expected relationship between workout length and energy expenditure.

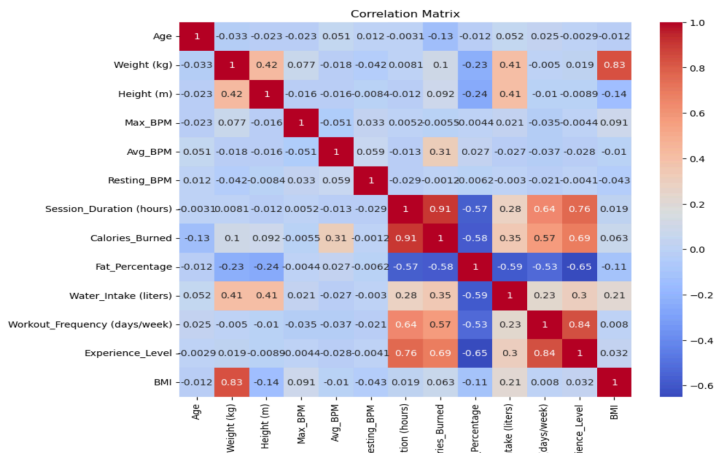

6. **Bivariate Analysis of Workout Frequency and Fitness Metrics**

   Bivariate analysis helped me explore relationships between two variables in more depth. For example:

   - **Water Intake vs. BMI**: A scatter plot was used to examine the relationship between water intake and BMI, revealing any trends between hydration and body composition.
   - **Age vs. Fat Percentage**: This analysis highlighted how age might correlate with body fat percentage, which can be crucial for tailoring workout recommendations based on age.
   - **Workout Frequency vs. Experience Level**: Using violin plots, I visualized workout frequency across different experience levels to observe whether experienced gym members worked out more consistently.

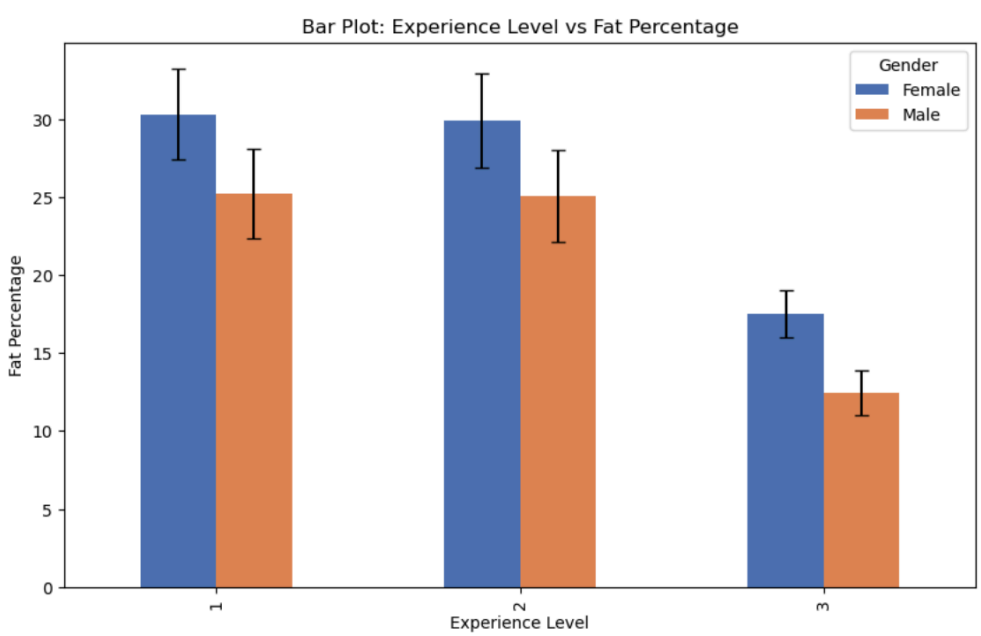

7. **Outlier Detection and Removal**
   Outliers can distort analysis, so I applied box plots across various metrics like session duration, calories burned, and BMI to detect any extreme values. Identifying and removing these outliers was essential to ensure that unusual data points didn't skew the findings. This step helped maintain a dataset that was representative of typical gym member behavior, making subsequent analyses and model predictions more reliable.
8. **Advanced EDA: Group Comparisons and Mean Values**
   To gain deeper insights, I also compared mean values across groups. For example, I calculated the average resting heart rate and calories burned for different age groups or experience levels. This comparison helped highlight specific characteristics, like whether certain age groups or experienced members tend to have lower resting heart rates, indicating better cardiovascular health.

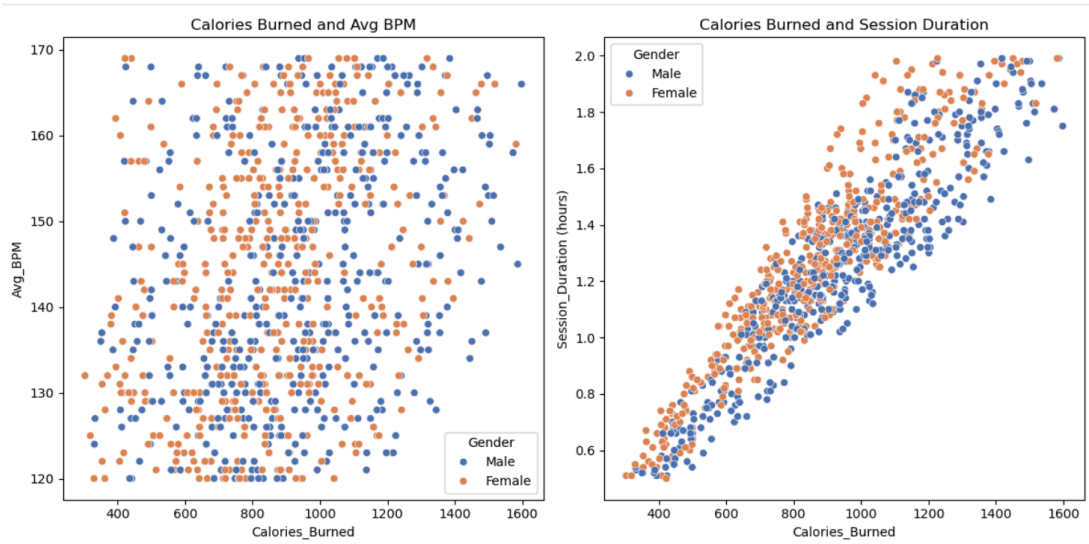

Through these analyses, I was able to derive comprehensive insights into gym members' workout habits and fitness levels. EDA not only provided context for understanding the data but also informed feature engineering and model-building decisions in later stages. These exploratory steps helped lay the groundwork for a more informed, data-driven approach to predicting caloric burn and understanding the factors that contribute to it.

Model Building and Selection: After completing the exploratory data analysis (EDA), I proceeded to build and evaluate several machine learning models to predict calorie burn based on gym members' demographic and workout data. Below, I break down each step:

1. **Feature Encoding and Scaling :** I started by preparing the data for modeling. Categorical variables like `Gender`, `Workout_Type`, and `Experience_Level` were transformed into numerical values using Label Encoding. This conversion allows models to interpret these features correctly. Next, I standardized the numerical features using Standard Scaling. Scaling ensures that each feature contributes equally to the model, preventing any single feature from dominating due to differing scales.

2. **Data Splitting :** To train and evaluate the models, I split the data into training and testing sets using an 80-20 split. This allows the model to learn from the training data and be evaluated on unseen test data, providing an indication of its generalization performance.
3. **Model Selection and Initialization :** I selected a range of regression algorithms to compare their effectiveness for this task:
    - **Linear Regression**
    - **Decision Tree Regressor**
    - **Random Forest Regressor**
    - **K-Neighbors Regressor**
    - **XGBoost Regressor**
4. Each model was initialized with specific hyperparameters to optimize performance. This range of models allows a comparison between simpler algorithms (like Linear Regression) and more complex ensemble methods (like Random Forest and XGBoost).
5. **Training and Evaluation :** Each model was trained on the training data and then evaluated on the test data. I used **Mean Squared Error (MSE)** and **R^2 Score** as evaluation metrics:
    - **Mean Squared Error** provides a measure of the average squared difference between actual and predicted calorie burn values. A lower MSE indicates better model performance.
    - **R^2 Score** represents the proportion of variance in the target variable that is predictable from the features. A higher R^2 score indicates a better fit.
6. Each model's performance metrics were printed and compared, helping identify which algorithm performed best for predicting calorie burn.
7. **Hyperparameter Tuning :** To further improve the model's performance, I employed **Grid Search Cross-Validation** on the Random Forest Regressor. Grid Search systematically tests combinations of hyperparameters to find the optimal set for the model. In this case, I tuned parameters like `n_estimators`, `max_depth`, and `min_samples_split` for the Random Forest model. The model achieved an optimal R^2 score with the best parameters found.

Model Evaluation and Comparison After hyperparameter tuning, I evaluated the models based on MSE and R2 Score. Here are the summarized results:

Model	**MSE**	**R2 Score**
Linear Regression	44731.43	0.38
Decision Tree Regressor	47805.11	0.34
Random Forest Regressor	31116.27	0.57
K-Neighbors Regressor	37801.84	0.47
XGBoost Regressor	33733.38	0.53

The Random Forest Regressor outperformed other models with the lowest MSE and highest R2 score, making it the best model for predicting caloric burn.

Insights and Observations: Through a comprehensive analysis of gym members' data, several key insights emerged that can guide strategic decisions for gyms. Firstly, strength training proved to be a popular choice among members, with a significant focus on muscle-building exercises. This indicates that marketing efforts should highlight muscle gains and strength-related goals to align with member interests. Additionally, the data reveals a slight gender disparity, with males making up a higher percentage of gym-goers. This suggests that targeted engagement strategies, such as female-only classes and inclusive marketing, could help attract more female members.

Cardio workouts are the most popular across both genders, presenting an opportunity to use cardio-focused programs as an entry point for diverse member engagement. Furthermore, middle-aged members (particularly those aged 46-55) form a substantial demographic in the gym and show an interest in both HIIT and strength training. Tailoring offerings to meet their needs, especially through age-appropriate strength-building programs, could enhance retention in this group. Additionally, younger members (18-25) show lower engagement, indicating a need for programs that align with their preferences, such as high-energy group workouts.

High body fat levels are notable within the gym population, with many members having high fat percentages, especially among older age groups. This highlights the potential for targeted fat reduction initiatives, such as weight management programs and specialized classes. Overall, the use of machine learning models, particularly the Random Forest Regressor, allowed for accurate calorie burn predictions, supporting personalized fitness plans and progress tracking. These insights underscore the value of data-driven strategies in optimizing gym services and meeting member needs effectively.

Conclusion: In conclusion, this analysis of gym members' data underscores the power of data science in uncovering insights that can enhance fitness programs and member engagement. By exploring demographic and behavioral trends, we can identify valuable patterns, such as the popularity of strength training, the slight male dominance among gym-goers, and the high engagement of middle-aged members. The analysis also highlighted opportunities for targeted fat reduction initiatives and the potential for increased participation among younger members through tailored programs. Using machine learning models, particularly the Random Forest Regressor, proved effective for predicting calorie burn, enabling personalized fitness recommendations. These insights provide a data-driven foundation for gyms to tailor their offerings, promote member retention, and create a more inclusive and engaging environment. Ultimately, this approach demonstrates how gyms can better meet the diverse fitness needs of their members by embracing data-driven strategies.